

```
In [1]: 1 import warnings
        2 warnings.filterwarnings('ignore')
```

```
In [2]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
```

## Loading Data

```
In [3]: 1 #Loading Data into data frame
        2 path = "https://s3-api.us-gio.objectstorage.softlayer.net/cf-courses-data/Cog
        3 cars = pd.read_csv(path,header = None)
```

```
In [4]: 1 #first five rows of the Data
        2 cars.head()
```

```
Out[4]:
```

	0	1	2	3	4	5	6	7	8	9	...	16	17	18	19	20	21
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	9.0	154
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.40	10.0	102
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	3.40	8.0	115

5 rows × 26 columns



```
In [5]: 1 #last five rows of the Data
        2 cars.tail()
```

```
Out[5]:
```

	0	1	2	3	4	5	6	7	8	9	...	16	17	18	19	20	21
200	-1	95	volvo	gas	std	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	9.5	114
201	-1	95	volvo	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	8.7	160
202	-1	95	volvo	gas	std	four	sedan	rwd	front	109.1	...	173	mpfi	3.58	2.87	8.8	134
203	-1	95	volvo	diesel	turbo	four	sedan	rwd	front	109.1	...	145	idi	3.01	3.40	23.0	106
204	-1	95	volvo	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	9.5	114

5 rows × 26 columns



## Add Headers

Take a look at our dataset; pandas automatically set the header by an integer from 0.

To better describe our data we can introduce a header, this information is available at:

<https://archive.ics.uci.edu/ml/datasets/Automobile>  
(<https://archive.ics.uci.edu/ml/datasets/Automobile>)

Thus, we have to add headers manually.

Firstly, we create a list "headers" that include all column names in order. Then, we use `dataframe.columns = headers` to replace the headers by the list we created.

```
In [6]: 1 # create headers list
2 headers = ["symboling", "normalized-losses", "make", "fuel-type", "aspiration", "
3           "drive-wheels", "engine-location", "wheel-base", "length", "width", "height",
4           "num-of-cylinders", "engine-size", "fuel-system", "bore", "stroke", "compression-ratio",
5           "peak-rpm", "city-mpg", "highway-mpg", "price"]
6 print("headers\n", headers)
```

headers

```
['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration', 'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke', 'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg', 'highway-mpg', 'price']
```

```
In [7]: 1 #assign headers to columns
2 cars.columns = headers
3
```

```
In [8]: 1 cars.head()
```

Out[8]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	.
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	.
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	.

5 rows × 26 columns



we can drop missing values along the column "price" as follows

In [9]: 1 cars.dropna(subset=["price"], axis=0)

Out[9]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	w
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	
3	2	164	audi	gas	std	four	sedan	fwd	front	
4	2	164	audi	gas	std	four	sedan	4wd	front	
5	2	?	audi	gas	std	two	sedan	fwd	front	
6	1	158	audi	gas	std	four	sedan	fwd	front	'
7	1	?	audi	gas	std	four	wagon	fwd	front	'
8	1	158	audi	gas	turbo	four	sedan	fwd	front	'
9	0	?	audi	gas	turbo	two	hatchback	4wd	front	
10	2	192	bmw	gas	std	two	sedan	rwd	front	'
11	0	192	bmw	gas	std	four	sedan	rwd	front	'
12	0	188	bmw	gas	std	two	sedan	rwd	front	'
13	0	188	bmw	gas	std	four	sedan	rwd	front	'
14	1	?	bmw	gas	std	four	sedan	rwd	front	'
15	0	?	bmw	gas	std	four	sedan	rwd	front	'
16	0	?	bmw	gas	std	two	sedan	rwd	front	'
17	0	?	bmw	gas	std	four	sedan	rwd	front	'
18	2	121	chevrolet	gas	std	two	hatchback	fwd	front	
19	1	98	chevrolet	gas	std	two	hatchback	fwd	front	
20	0	81	chevrolet	gas	std	four	sedan	fwd	front	
21	1	118	dodge	gas	std	two	hatchback	fwd	front	
22	1	118	dodge	gas	std	two	hatchback	fwd	front	
23	1	118	dodge	gas	turbo	two	hatchback	fwd	front	
24	1	148	dodge	gas	std	four	hatchback	fwd	front	
25	1	148	dodge	gas	std	four	sedan	fwd	front	
26	1	148	dodge	gas	std	four	sedan	fwd	front	
27	1	148	dodge	gas	turbo	?	sedan	fwd	front	
28	-1	110	dodge	gas	std	four	wagon	fwd	front	'
29	3	145	dodge	gas	turbo	two	hatchback	fwd	front	
...	...	...	...	...	...	...	...	...	...	
175	-1	65	toyota	gas	std	four	hatchback	fwd	front	'
176	-1	65	toyota	gas	std	four	sedan	fwd	front	'

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	w
177	-1	65	toyota	gas	std	four	hatchback	fwd	front	'
178	3	197	toyota	gas	std	two	hatchback	rwd	front	'
179	3	197	toyota	gas	std	two	hatchback	rwd	front	'
180	-1	90	toyota	gas	std	four	sedan	rwd	front	'
181	-1	?	toyota	gas	std	four	wagon	rwd	front	'
182	2	122	volkswagen	diesel	std	two	sedan	fwd	front	
183	2	122	volkswagen	gas	std	two	sedan	fwd	front	
184	2	94	volkswagen	diesel	std	four	sedan	fwd	front	
185	2	94	volkswagen	gas	std	four	sedan	fwd	front	
186	2	94	volkswagen	gas	std	four	sedan	fwd	front	
187	2	94	volkswagen	diesel	turbo	four	sedan	fwd	front	
188	2	94	volkswagen	gas	std	four	sedan	fwd	front	
189	3	?	volkswagen	gas	std	two	convertible	fwd	front	
190	3	256	volkswagen	gas	std	two	hatchback	fwd	front	
191	0	?	volkswagen	gas	std	four	sedan	fwd	front	'
192	0	?	volkswagen	diesel	turbo	four	sedan	fwd	front	'
193	0	?	volkswagen	gas	std	four	wagon	fwd	front	'
194	-2	103	volvo	gas	std	four	sedan	rwd	front	'
195	-1	74	volvo	gas	std	four	wagon	rwd	front	'
196	-2	103	volvo	gas	std	four	sedan	rwd	front	'
197	-1	74	volvo	gas	std	four	wagon	rwd	front	'
198	-2	103	volvo	gas	turbo	four	sedan	rwd	front	'
199	-1	74	volvo	gas	turbo	four	wagon	rwd	front	'
200	-1	95	volvo	gas	std	four	sedan	rwd	front	'
201	-1	95	volvo	gas	turbo	four	sedan	rwd	front	'
202	-1	95	volvo	gas	std	four	sedan	rwd	front	'
203	-1	95	volvo	diesel	turbo	four	sedan	rwd	front	'
204	-1	95	volvo	gas	turbo	four	sedan	rwd	front	'

205 rows × 26 columns



Now, we have successfully read the raw dataset and add the correct headers into the data frame.

```
In [10]: 1 cars.columns
```

```
Out[10]: Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
               'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
               'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
               'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
               'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
               'highway-mpg', 'price'],
              dtype='object')
```

```
In [11]: 1 #to save the file in local directory with suitable columns.
         2 cars.to_csv("cars.csv",index=False)
```

```
In [12]: 1 cars.shape
```

```
Out[12]: (205, 26)
```

## Data Wrangling

```
In [13]: 1 #check the first five rows of data:
         2 cars.head()
```

```
Out[13]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4

5 rows × 26 columns



As we can see, several question marks appeared in the dataframe; those are missing values which may hinder our further analysis.

So, how do we identify all those missing values and deal with them?

### How to work with missing data?

Steps for working with missing data:

1. identify missing data
2. deal with missing data
3. correct data format

# Identify and handle missing values

## Identify missing values

### Convert "?" to NaN

In the car dataset, missing data comes with the question mark "?". We replace "?" with NaN (Not a Number), which is Python's default missing value marker, for reasons of computational speed and convenience. Here we use the function:

```
.replace(A, B, inplace = True)
```

to replace A by B

```
In [14]: 1 # replace "?" to NaN
          2 cars.replace("?", np.nan, inplace = True)
          3 cars.head(5)
```

```
Out[14]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	
0	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
1	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
2	1	NaN	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	.
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	.
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	.

5 rows × 26 columns

```
In [15]: 1 missing_data = cars.isnull()
          2 missing_data.head(5)
```

```
Out[15]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	er
0	False	True	False	False	False	False	False	False	False	False	...	
1	False	True	False	False	False	False	False	False	False	False	...	
2	False	True	False	False	False	False	False	False	False	False	...	
3	False	False	False	False	False	False	False	False	False	False	...	
4	False	False	False	False	False	False	False	False	False	False	...	

5 rows × 26 columns

## Count missing values in each column

Using a for loop in Python, we can quickly figure out the number of missing values in each column. As mentioned above, "True" represents a missing value, "False" means the value is present in the dataset. In the body of the for loop the method ".value\_counts()" counts the number of "True" values.

```
In [16]: 1 for column in missing_data.columns.values.tolist():
          2     print(column)
          3     print (missing_data[column].value_counts())
          4     print("")
```

```
symboling
False      205
Name: symboling, dtype: int64

normalized-losses
False      164
True        41
Name: normalized-losses, dtype: int64

make
False      205
Name: make, dtype: int64

fuel-type
False      205
Name: fuel-type, dtype: int64

aspiration
False      205
Name: aspiration, dtype: int64
```

Based on the summary above, each column has 205 rows of data, seven columns containing missing data:

1. "normalized-losses": 41 missing data
2. "num-of-doors": 2 missing data
3. "bore": 4 missing data
4. "stroke" : 4 missing data
5. "horsepower": 2 missing data
6. "peak-rpm": 2 missing data
7. "price": 4 missing data

## Deal with missing data

### How to deal with missing data?

1. drop data
  - a. drop the whole row
  - b. drop the whole column
2. replace data
  - a. replace it by mean

- b. replace it by frequency
- c. replace it based on other functions

Whole columns should be dropped only if most entries in the column are empty. In our dataset, none of the columns are empty enough to drop entirely. We have some freedom in choosing which method to replace data; however, some methods may seem more reasonable than others. We will apply each method to many different columns:

#### Replace by mean:

- "normalized-losses": 41 missing data, replace them with mean
- "stroke": 4 missing data, replace them with mean
- "bore": 4 missing data, replace them with mean
- "horsepower": 2 missing data, replace them with mean
- "peak-rpm": 2 missing data, replace them with mean

#### Replace by frequency:

- "num-of-doors": 2 missing data, replace them with "four".
  - Reason: 84% sedans is four doors. Since four doors is most frequent, it is most likely to occur

#### Drop the whole row:

- "price": 4 missing data, simply delete the whole row
  - Reason: price is what we want to predict. Any data entry without price data cannot be used for prediction; therefore any row now without price data is not useful to us

#### Calculate the average of the column "normalized-losses"

```
In [17]: 1 avg_norm_loss = cars["normalized-losses"].astype("float").mean(axis=0)
          2 print("Average of normalized-losses:", avg_norm_loss)
```

Average of normalized-losses: 122.0

#### Replace "NaN" by mean value in "normalized-losses" column

```
In [18]: 1 cars["normalized-losses"].replace(np.nan, avg_norm_loss, inplace=True)
```

```
In [19]: 1 cars["normalized-losses"].head()
```

```
Out[19]: 0    122
          1    122
          2    122
          3    164
          4    164
          Name: normalized-losses, dtype: object
```

Now the missing values in the column "normalized-losses" has been replaced by the mean



Now the missing values in the column "normalized-residuals" has been replaced by the mean of the column.

### Dealing with the missing value in "num-of-doors"

```
In [20]: 1 cars["num-of-doors"].value_counts()  
        2
```

```
Out[20]: four      114  
        two       89  
        Name: num-of-doors, dtype: int64
```

```
In [21]: 1 #maximum freq. in num-of-doors  
        2 cars["num-of-doors"].value_counts().idxmax()
```

```
Out[21]: 'four'
```

***It is clear that four is the most frequent num-of-doors in cars so replace the missing value with the four in the "num-of-doors" column***

```
In [22]: 1 cars["num-of-doors"].replace(np.nan, "four", inplace=True)
```

```
In [23]: 1 cars["num-of-doors"].value_counts()
```

```
Out[23]: four      116  
        two       89  
        Name: num-of-doors, dtype: int64
```

### Dealing with the missing value in "bore"

```
In [24]: 1 cars["bore"].head()
```

```
Out[24]: 0      3.47  
        1      3.47  
        2      2.68  
        3      3.19  
        4      3.19  
        Name: bore, dtype: object
```

### Calculate the mean value for 'bore' column

```
In [25]: 1 avg_bore=cars['bore'].astype('float').mean(axis=0)  
        2 print("Average of bore:", avg_bore)
```

```
Average of bore: 3.3297512437810957
```

### Replace NaN by mean value

```
In [26]: 1 cars["bore"].replace(np.nan, avg_bore, inplace=True)
```

### Dealing with the missing value in "stroke"

```
In [27]: 1 cars["stroke"].head()
```

```
Out[27]: 0    2.68
         1    2.68
         2    3.47
         3    3.40
         4    3.40
         Name: stroke, dtype: object
```

### Calculate the mean value for 'bore' column

```
In [28]: 1 avg_stroke=cars['stroke'].astype('float').mean(axis=0)
         2 print("Average of stroke:", avg_stroke)
```

Average of stroke: 3.2554228855721337

### Replace NaN by mean value

```
In [29]: 1 cars["stroke"].replace(np.nan, avg_bore, inplace=True)
```

### Dealing with the missing value in "horsepower"

```
In [30]: 1 cars["horsepower"].head()
```

```
Out[30]: 0    111
         1    111
         2    154
         3    102
         4    115
         Name: horsepower, dtype: object
```

### Calculate the mean value for 'horsepower' column

```
In [31]: 1 avg_horsepower=cars['horsepower'].astype('float').mean(axis=0)
         2 print("Average of horsepower:", avg_horsepower)
```

Average of horsepower: 104.25615763546799

### Replace NaN by mean value

```
In [32]: 1 cars["horsepower"].replace(np.nan, avg_horsepower, inplace=True)
```

### Dealing with the missing value in "peak-rpm"

```
In [33]: 1 cars["peak-rpm"].head()
```

```
Out[33]: 0    5000
         1    5000
         2    5000
         3    5500
         4    5500
         Name: peak-rpm, dtype: object
```

### Calculate the mean value for 'peak-rpm' column

```
In [34]: 1 avg_peak_rpm = cars["peak-rpm"].astype('float').mean(axis=0)
         2 print("Average of peak-rpm:", avg_peak_rpm)
         3
```

Average of peak-rpm: 5125.369458128079

### Replace NaN by mean value

```
In [35]: 1 cars["peak-rpm"].replace(np.nan, avg_peak_rpm, inplace=True)
```

### Dealing with the missing value in "price"

```
In [36]: 1 cars["price"].head()
```

```
Out[36]: 0    13495
         1    16500
         2    16500
         3    13950
         4    17450
         Name: price, dtype: object
```

Since the price is what we want to predict so any empty entry is of no use. Hence, Drop the entire row which contain null value.

```
In [37]: 1 # simply drop whole row with NaN in "price" column
         2 cars.dropna(subset=["price"], axis=0, inplace=True)
         3
         4 # reset index, because we dropped two rows
         5 cars.reset_index(drop=True, inplace=True)
```

In [47]: 1 cars.head()

Out[47]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	.
0	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
1	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
2	1	122	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	.
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	.
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	.

5 rows × 26 columns



Now, we obtain the dataset with no missing values.

## Correct data format

In Pandas, we use

**.dtype()** to check the data type

**.astype()** to change the data type

**Lets list the data types for each column**

```
In [38]: 1 cars.dtypes
```

```
Out[38]: symboling          int64
normalized-losses         object
make                      object
fuel-type                 object
aspiration                object
num-of-doors              object
body-style                object
drive-wheels              object
engine-location           object
wheel-base                float64
length                   float64
width                     float64
height                    float64
curb-weight               int64
engine-type               object
num-of-cylinders           object
engine-size               int64
fuel-system               object
bore                      object
stroke                    object
compression-ratio         float64
horsepower                 object
peak-rpm                  object
city-mpg                   int64
highway-mpg                int64
price                     object
dtype: object
```

As we can see above, some columns are not of the correct data type. Numerical variables should have type 'float' or 'int', and variables with strings such as categories should have type 'object'. For example, 'bore' and 'stroke' variables are numerical values that describe the engines, so we should expect them to be of the type 'float' or 'int'; however, they are shown as type 'object'. We have to convert data types into a proper format for each column.

### Convert data types to proper format

```
In [39]: 1 cars[["bore", "stroke"]] = cars[["bore", "stroke"]].astype("float")
2 cars[["normalized-losses"]] = cars[["normalized-losses"]].astype("int")
3 cars[["price"]] = cars[["price"]].astype("float")
4 cars[["peak-rpm"]] = cars[["peak-rpm"]].astype("float")
```

Let us list the columns after the conversion

```
In [40]: 1 cars.dtypes
```

```
Out[40]: symboling          int64
normalized-losses      int32
make                   object
fuel-type              object
aspiration             object
num-of-doors          object
body-style            object
drive-wheels          object
engine-location        object
wheel-base            float64
length                float64
width                  float64
height                float64
curb-weight            int64
engine-type            object
num-of-cylinders       object
engine-size            int64
fuel-system            object
bore                   float64
stroke                float64
compression-ratio      float64
horsepower             object
peak-rpm               float64
city-mpg               int64
highway-mpg            int64
price                  float64
dtype: object
```

Now, we finally obtain the cleaned dataset with no missing values and all data in its proper format.

## Data Standardization

Data is usually collected from different agencies with different formats. (Data Standardization is also a term for a particular type of data normalization, where we subtract the mean and divide by the standard deviation)

### Example

Transform mpg to L/100km:

In our dataset, the fuel consumption columns "city-mpg" and "highway-mpg" are represented by mpg (miles per gallon) unit. Assume we are developing an application in a country that accept the fuel consumption with L/100km standard

We will need to apply **data transformation** to transform mpg into L/100km?

The formula for unit conversion is

$$\text{L/100km} = 235 / \text{mpg}$$

We can do many mathematical operations directly in Pandas.

In [41]:

```
1 cars.head()
```

Out[41]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	
0	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
1	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
2	1	122	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	.
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	.
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	.

5 rows × 26 columns



In [42]:

```
1 # Convert mpg to L/100km by mathematical operation (235 divided by mpg)
2 cars['city-L/100km'] = 235/cars["city-mpg"]
3
4 # check transformed data
5 cars.head()
```

Out[42]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	
0	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
1	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
2	1	122	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	.
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	.
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	.

5 rows × 27 columns



In [43]:

```
1 #now drop the column("city-mpg") as we already created a standard version as
2 cars.drop(columns=['city-mpg'],inplace=True)
```

In [44]: 1 cars.head()

Out[44]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	
0	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
1	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
2	1	122	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	.
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	.
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	.

5 rows × 26 columns



```

In [45]: 1 # transform mpg to L/100km by mathematical operation (235 divided by mpg)
2 cars["highway-mpg"] = 235/cars["highway-mpg"]
3
4 # rename column name from "highway-mpg" to "highway-L/100km"
5 cars.rename(columns={"highway-mpg": "highway-L/100km"}, inplace= True)
6
7 # check your transformed data
8 cars.head()
9

```

Out[45]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	
0	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
1	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
2	1	122	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	.
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	.
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	.

5 rows × 26 columns



## Data Normalization

let's say we want to scale the columns "length", "width" and "height"

**Target:**would like to Normalize those variables so their value ranges from 0 to 1.



**Approach:** replace original value by (original value)/(maximum value)

In [47]: `1 cars[["length", "width", "height"]].head()`

Out[47]:

	length	width	height
0	168.8	64.1	48.8
1	168.8	64.1	48.8
2	171.2	65.5	52.4
3	176.6	66.2	54.3
4	176.6	66.4	54.3

In [48]: `1 # replace (original value) by (original value)/(maximum value)  
2 cars['length'] = cars['length']/cars['length'].max()  
3 cars['width'] = cars['width']/cars['width'].max()  
4 cars['height'] = cars['height']/cars['height'].max()  
5 # show the scaled columns  
6 cars[["length", "width", "height"]].head()`

Out[48]:

	length	width	height
0	0.811148	0.890278	0.816054
1	0.811148	0.890278	0.816054
2	0.822681	0.909722	0.876254
3	0.848630	0.919444	0.908027
4	0.848630	0.922222	0.908027

## Binning

In our dataset, "horsepower" is a real valued variable ranging from 48 to 288, it has 57 unique values. What if we only care about the price difference between cars with high horsepower, medium horsepower, and little horsepower (3 types)? Can we rearrange them into three 'bins' to simplify analysis?

We will use the Pandas method 'cut' to segment the 'horsepower' column into 3 bins

Convert data to correct format

In [49]: `1 cars["horsepower"] = cars["horsepower"].astype(int, copy=True)`


In [50]: `1 cars["horsepower"].dtypes`

Out[50]: `dtype('int32')`

Lets plot the histogram of horsepower, to see what the distribution of horsepower looks like.

```
In [51]: 1 %matplotlib inline
2 import matplotlib as plt
3 from matplotlib import pyplot
4 plt.pyplot.hist(cars["horsepower"])
5
6 # set x/y labels and plot title
7 plt.pyplot.xlabel("horsepower")
8 plt.pyplot.ylabel("count")
9 plt.pyplot.title("horsepower bins")
```

Out[51]: Text(0.5, 1.0, 'horsepower bins')



We would like 3 bins of equal size bandwidth so we use numpy's `linspace(start_value, end_value, numbers_generated)` function.

Since we want to include the minimum value of horsepower we want to set `start_value=min(df["horsepower"])`.

Since we want to include the maximum value of horsepower we want to set `end_value=max(df["horsepower"])`.

Since we are building 3 bins of equal length, there should be 4 dividers, so `numbers_generated=4`.

We build a bin array, with a minimum value to a maximum value, with bandwidth calculated above. The bins will be values used to determine when one bin ends and another begins.

```
In [52]: 1 bins = np.linspace(min(cars["horsepower"]), max(cars["horsepower"]), 4)
2 bins
```

Out[52]: array([ 48. , 119.33333333, 190.66666667, 262. ])

We set group names:

```
In [53]: 1 group_names = ['Low', 'Medium', 'High']
```

We apply the function "cut" to determine what each value of "df['horsepower']" belongs to.

```
In [54]: 1 cars['horsepower-binned'] = pd.cut(cars['horsepower'], bins, labels=group_names)
2 cars[['horsepower', 'horsepower-binned']].head(20)
```

```
Out[54]:
```

	horsepower	horsepower-binned
0	111	Low
1	111	Low
2	154	Medium
3	102	Low
4	115	Low
5	110	Low
6	110	Low
7	110	Low
8	140	Medium
9	101	Low
10	101	Low
11	121	Medium
12	121	Medium
13	121	Medium
14	182	Medium
15	182	Medium
16	182	Medium
17	48	Low
18	70	Low
19	70	Low

Lets see the number of vehicles in each bin.


```
In [55]: 1 cars["horsepower-binned"].value_counts()
```

```
Out[55]: Low      153
Medium    43
High       5
Name: horsepower-binned, dtype: int64
```

Lets plot the distribution of each bin.

```
In [56]: 1 %matplotlib inline
2 import matplotlib as plt
3 from matplotlib import pyplot
4 pyplot.bar(group_names, cars["horsepower-binned"].value_counts())
5
6 # set x/y labels and plot title
7 plt.pyplot.xlabel("horsepower")
8 plt.pyplot.ylabel("count")
9 plt.pyplot.title("horsepower bins")
```

Out[56]: Text(0.5, 1.0, 'horsepower bins')




```
In [57]: 1 cars.head()
```

Out[57]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base
0	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6
1	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6
2	1	122	alfa-romero	gas	std	two	hatchback	rwd	front	94.5
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4

5 rows × 27 columns




## Bins visualization

Normally, a histogram is used to visualize the distribution of bins we created above.

```
In [58]: 1 %matplotlib inline
2 import matplotlib as plt
3 from matplotlib import pyplot
4
5
6 # draw histogram of attribute "horsepower" with bins = 3
7 plt.pyplot.hist(cars["horsepower"], bins = 3)
8
9 # set x/y labels and plot title
10 plt.pyplot.xlabel("horsepower")
11 plt.pyplot.ylabel("count")
12 plt.pyplot.title("horsepower bins")
```

Out[58]: Text(0.5, 1.0, 'horsepower bins')



The plot above shows the binning result for attribute "horsepower".

## Indicator variable (or dummy variable)

We see the column "fuel-type" has two unique values, "gas" or "diesel". Regression doesn't understand words, only numbers. To use this attribute in regression analysis, we convert "fuel-type" into indicator variables.

We will use the panda's method 'get\_dummies' to assign numerical values to different categories of fuel type.

```
In [59]: 1 cars.columns
```

```
Out[59]: Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
               'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
               'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
               'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
               'compression-ratio', 'horsepower', 'peak-rpm', 'highway-L/100km',
               'price', 'city-L/100km', 'horsepower-binned'],
              dtype='object')
```

get indicator variables and assign it to data frame "dummy\_variable\_1"

```
In [60]: 1 dummy_variable_1 = pd.get_dummies(cars["fuel-type"])
        2 dummy_variable_1.head()
```

```
Out[60]:
```

	diesel	gas
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1


We now have the value 0 to represent "gas" and 1 to represent "diesel" in the column "fuel-type". We will now insert this column back into our original dataset.

```
In [61]: 1 cars.head()
        2
```

```
Out[61]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base
0	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6
1	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6
2	1	122	alfa-romero	gas	std	two	hatchback	rwd	front	94.5
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4

5 rows × 11 columns



```
In [62]: 1 # merge data frame "cars" and "dummy_variable_1"
        2 cars = pd.concat([cars, dummy_variable_1], axis=1)
        3
        4 # drop original column "fuel-type" from "cars"
        5 cars.drop("fuel-type", axis = 1, inplace=True)
```


In [63]:

```
1 cars.head()
```

Out[63]:

	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length
0	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.81114
1	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.81114
2	1	122	alfa-romero	std	two	hatchback	rwd	front	94.5	0.82268
3	2	164	audi	std	four	sedan	fwd	front	99.8	0.84863
4	2	164	audi	std	four	sedan	4wd	front	99.4	0.84863

5 rows × 28 columns



The last two columns are now the indicator variable representation of the fuel-type variable. It's all 0s and 1s now.


In [64]:

```
1 cars.rename(columns={'gas':'fuel-type-gas', 'diesel':'fuel-type-diesel'}, inplace=True)
2 cars.head()
```

Out[64]:

	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length
0	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.81114
1	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.81114
2	1	122	alfa-romero	std	two	hatchback	rwd	front	94.5	0.82268
3	2	164	audi	std	four	sedan	fwd	front	99.8	0.84863
4	2	164	audi	std	four	sedan	4wd	front	99.4	0.84863

5 rows × 28 columns



```
In [65]: 1 # get indicator variables of aspiration and assign it to data frame "dummy_va
2 dummy_variable_2 = pd.get_dummies(cars['aspiration'])
3
4 # change column names for clarity
5 dummy_variable_2.rename(columns={'std':'aspiration-std', 'turbo': 'aspiration
6
7 # show first 5 instances of data frame "dummy_variable_1"
8 dummy_variable_2.head()
```

```
Out[65]:
```


	aspiration-std	aspiration-turbo
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0

```
In [66]: 1
2 #merge the new dataframe to the original dataframe
3 cars = pd.concat([cars, dummy_variable_2], axis=1)
4
5 cars.head()
```

```
Out[66]:
```

	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length
0	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.81114
1	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.81114
2	1	122	alfa-romero	std	two	hatchback	rwd	front	94.5	0.82268
3	2	164	audi	std	four	sedan	fwd	front	99.8	0.84863
4	2	164	audi	std	four	sedan	4wd	front	99.4	0.84863

5 rows × 30 columns






```
In [67]: 1 # drop original column "aspiration" from "df"
2 cars.drop('aspiration', axis = 1, inplace=True)
3 cars.head()
```

Out[67]:

	symboling	normalized-losses	make	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	width
0	3	122	alfa-romero	two	convertible	rwd	front	88.6	0.811148	0.890278
1	3	122	alfa-romero	two	convertible	rwd	front	88.6	0.811148	0.890278
2	1	122	alfa-romero	two	hatchback	rwd	front	94.5	0.822681	0.909722
3	2	164	audi	four	sedan	fwd	front	99.8	0.848630	0.919444
4	2	164	audi	four	sedan	4wd	front	99.4	0.848630	0.922222

5 rows × 29 columns



save the new csv

```
In [68]: 1 cars.to_csv('clean_cars.csv')
```

## Exploratory-Data-Analysis

Analyzing Individual Feature Patterns using Visualization

To install seaborn we use the pip which is the python package manager.

```
In [69]: 1 %%capture
2 ! pip install seaborn
```

```
In [70]: 1 import matplotlib.pyplot as plt
2 import seaborn as sns
3 %matplotlib inline
```

Finding Correlation

In [71]: 1 cars.corr()


Out[71]:

	symboling	normalized-losses	wheel-base	length	width	height	curb-weight	engine-size
<b>symboling</b>	1.000000	0.466264	-0.535987	-0.365404	-0.242423	-0.550160	-0.233118	-0.110
<b>normalized-losses</b>	0.466264	1.000000	-0.056661	0.019424	0.086802	-0.373737	0.099404	0.112
<b>wheel-base</b>	-0.535987	-0.056661	1.000000	0.876024	0.814507	0.590742	0.782097	0.572
<b>length</b>	-0.365404	0.019424	0.876024	1.000000	0.857170	0.492063	0.880665	0.685
<b>width</b>	-0.242423	0.086802	0.814507	0.857170	1.000000	0.306002	0.866201	0.729
<b>height</b>	-0.550160	-0.373737	0.590742	0.492063	0.306002	1.000000	0.307581	0.074
<b>curb-weight</b>	-0.233118	0.099404	0.782097	0.880665	0.866201	0.307581	1.000000	0.849
<b>engine-size</b>	-0.110581	0.112360	0.572027	0.685025	0.729436	0.074694	0.849072	1.000
<b>bore</b>	-0.140019	-0.029862	0.493244	0.608971	0.544885	0.180449	0.644060	0.572
<b>stroke</b>	-0.000059	0.059131	0.155225	0.121904	0.188301	-0.068633	0.166038	0.199
<b>compression-ratio</b>	-0.182196	-0.114713	0.250313	0.159733	0.189867	0.259737	0.156433	0.028
<b>horsepower</b>	0.075810	0.217300	0.371178	0.579795	0.615056	-0.087001	0.757981	0.822
<b>peak-rpm</b>	0.279740	0.239543	-0.360305	-0.285970	-0.245800	-0.309974	-0.279361	-0.256
<b>highway-L/100km</b>	-0.029807	0.181189	0.577576	0.707108	0.736728	0.084301	0.836921	0.783
<b>price</b>	-0.082391	0.133999	0.584642	0.690628	0.751265	0.135486	0.834415	0.872
<b>city-L/100km</b>	0.066171	0.238567	0.476153	0.657373	0.673363	0.003811	0.785353	0.745
<b>fuel-type-diesel</b>	-0.196735	-0.101546	0.307237	0.211187	0.244356	0.281578	0.221046	0.070
<b>fuel-type-gas</b>	0.196735	0.101546	-0.307237	-0.211187	-0.244356	-0.281578	-0.221046	-0.070
<b>aspiration-std</b>	0.054615	0.006911	-0.256889	-0.230085	-0.305732	-0.090336	-0.321955	-0.110
<b>aspiration-turbo</b>	-0.054615	-0.006911	0.256889	0.230085	0.305732	0.090336	0.321955	0.110

Let's find the scatterplot of "engine-size" and "price"

```
In [72]: 1 # Engine size as potential predictor variable of price
          2 sns.regplot(x="engine-size", y="price", data=cars)
          3 plt.ylim(0,)
```

Out[72]: (0, 56269.964617255086)



As the engine-size goes up, the price goes up: this indicates a positive direct correlation between these two variables. Engine size seems like a pretty good predictor of price since the regression line is almost a perfect diagonal line.

We can examine the correlation between 'engine-size' and 'price'

```
In [73]: 1 cars[["engine-size", "price"]].corr()
```

```
Out[73]:
```


	engine-size	price
engine-size	1.000000	0.872335
price	0.872335	1.000000

it's approximately 0.87

Highway mpg is a potential predictor variable of price

```
In [74]: 1 sns.regplot(x="highway-L/100km", y="price", data=cars)
```

```
Out[74]: <matplotlib.axes._subplots.AxesSubplot at 0x2172b8f3dd8>
```



As the highway-L/100km goes up, the price goes up: this indicates an positive relationship between these two variables. highway-L/100km could potentially be a predictor of price.

We can examine the correlation between 'highway-mpg' and 'price'

```
In [75]: 1 cars[['highway-L/100km', 'price']].corr()
```

```
Out[75]:
```


	highway-L/100km	price
highway-L/100km	1.000000	0.801118
price	0.801118	1.000000

it's approximately 0.801118

Let's see if "Peak-rpm" as a predictor variable of "price".

```
In [76]: 1 sns.regplot(x="peak-rpm", y="price", data=cars)
```

```
Out[76]: <matplotlib.axes._subplots.AxesSubplot at 0x2172b9601d0>
```



Peak rpm does not seem like a good predictor of the price at all since the regression line is close to horizontal. Also, the data points are very scattered and far from the fitted line, showing lots of variability. Therefore it's not a reliable variable.

We can examine the correlation between 'peak-rpm' and 'price'

```
In [77]: 1 cars[['peak-rpm', 'price']].corr()
```

```
Out[77]:
```

	peak-rpm	price
peak-rpm	1.000000	-0.101616
price	-0.101616	1.000000

it's approximately -0.101616

Let's see if "stroke" as a predictor variable of "price".


```
In [78]: 1 cars[['stroke','price']].corr()
```

```
Out[78]:
```

	stroke	price
stroke	1.000000	0.082116
price	0.082116	1.000000

```
In [79]: 1 sns.regplot(x='stroke',y='price',data=cars)
```

```
Out[79]: <matplotlib.axes._subplots.AxesSubplot at 0x2172b9ce2b0>
```



There is a weak correlation between the variable 'stroke' and 'price.' as such regression will not work well.


## Categorical variables

These are variables that describe a 'characteristic' of a data unit, and are selected from a small group of categories. The categorical variables can have the type "object" or "int64". A good way to visualize categorical variables is by using boxplots.

Let's look at the relationship between "body-style" and "price".

```
In [80]: 1 sns.boxplot(x="body-style", y="price", data=cars)
```


```
Out[80]: <matplotlib.axes._subplots.AxesSubplot at 0x2172ba110f0>
```



We see that the distributions of price between the different body-style categories have a significant overlap, and so body-style would not be a good predictor of price. Let's examine engine "engine-location" and "price":

```
In [81]: 1 sns.boxplot(x="engine-location", y="price", data=cars)
```

```
Out[81]: <matplotlib.axes._subplots.AxesSubplot at 0x2172ba91080>
```




Here we see that the distribution of price between these two engine-location categories, front and rear, are distinct enough to take engine-location as a potential good predictor of price.

Let's examine "drive-wheels" and "price".



```
In [82]: 1 # drive-wheels  
2 sns.boxplot(x="drive-wheels", y="price", data=cars)
```

```
Out[82]: <matplotlib.axes._subplots.AxesSubplot at 0x2172bb0bda0>
```



Here we see that the distribution of price between the different drive-wheels categories differs; as such drive-wheels could potentially be a predictor of price.

## Descriptive Statistical Analysis

Let's first take a look at the variables by utilizing a description method.

The **describe** function automatically computes basic statistics for all continuous variables. Any NaN values are automatically skipped in these statistics.

This will show:

- the count of that variable
- the mean
- the standard deviation (std)
- the minimum value
- the IQR (Interquartile Range: 25%, 50% and 75%)
- the maximum value

In [83]: 1 cars.describe()

Out[83]:

	symboling	normalized-losses	wheel-base	length	width	height	curb-weight	e
count	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000
mean	0.840796	122.000000	98.797015	0.837102	0.915126	0.899108	2555.666667	126.8
std	1.254802	31.996250	6.066366	0.059213	0.029187	0.040933	517.296727	41.5
min	-2.000000	65.000000	86.600000	0.678039	0.837500	0.799331	1488.000000	61.0
25%	0.000000	101.000000	94.500000	0.801538	0.890278	0.869565	2169.000000	98.0
50%	1.000000	122.000000	97.000000	0.832292	0.909722	0.904682	2414.000000	120.0
75%	2.000000	137.000000	102.400000	0.881788	0.925000	0.928094	2926.000000	141.0
max	3.000000	256.000000	120.900000	1.000000	1.000000	1.000000	4066.000000	326.0

The default setting of "describe" skips variables of type object. We can apply the method "describe" on the variables of type 'object' as follows:

In [84]: 1 cars.describe(include=['object'])

Out[84]:

	make	num-of-doors	body-style	drive-wheels	engine-location	engine-type	num-of-cylinders	fuel-system
count	201	201	201	201	201	201	201	201
unique	22	2	5	3	2	6	7	8
top	toyota	four	sedan	fwd	front	ohc	four	mpfi
freq	32	115	94	118	198	145	157	92

## Value Counts

Value-counts is a good way of understanding how many units of each characteristic/variable we have. We can apply the "value\_counts" method on the column 'drive-wheels'. Don't forget the method "value\_counts" only works on Pandas series, not Pandas Dataframes. As a result, we only include one bracket "df['drive-wheels']" not two brackets "df[['drive-wheels']]".

In [85]: 1 cars['drive-wheels'].value\_counts()

Out[85]: fwd 118  
rwd 75  
4wd 8  
Name: drive-wheels, dtype: int64

We can convert the series to a Dataframe as follows :

In [86]: `1 cars['drive-wheels'].value_counts().to_frame()`

Out[86]:

drive-wheels	
fwd	118
rwd	75
4wd	8

Let's repeat the above steps but save the results to the dataframe "drive\_wheels\_counts" and rename the column 'drive-wheels' to 'value\_counts'.

In [87]: `1 drive_wheels_counts = cars['drive-wheels'].value_counts().to_frame()  
2 drive_wheels_counts.rename(columns={'drive-wheels': 'value_counts'}, inplace=True)  
3 drive_wheels_counts`

Out[87]:

value_counts	
fwd	118
rwd	75
4wd	8

Now let's rename the index to 'drive-wheels':

In [88]: `1 drive_wheels_counts.index.name = 'drive-wheels'  
2 drive_wheels_counts`

Out[88]:

value_counts	
drive-wheels	
fwd	118
rwd	75
4wd	8

We can repeat the above process for the variable 'engine-location'.

In [89]: `1 # engine-location as variable  
2 engine_loc_counts = cars['engine-location'].value_counts().to_frame()  
3 engine_loc_counts.rename(columns={'engine-location': 'value_counts'}, inplace=True)  
4 engine_loc_counts.index.name = 'engine-location'  
5 engine_loc_counts.head(10)`

Out[89]:

value_counts	
engine-location	
front	198
rear	3

Examining the value counts of the engine location would not be a good predictor variable for the

price. This is because we only have three cars with a rear engine and 198 with an engine in the front, this result is skewed. Thus, we are not able to draw any conclusions about the engine location.

## Basics of Grouping

The "groupby" method groups data by different categories. The data is grouped based on one or several variables and analysis is performed on the individual groups.

For example, let's group by the variable "drive-wheels". We see that there are 3 different categories of drive wheels.

```
In [90]: 1 cars['drive-wheels'].unique()
```

```
Out[90]: array(['rwd', 'fwd', '4wd'], dtype=object)
```

If we want to know, on average, which type of drive wheel is most valuable, we can group "drive-wheels" and then average them.

We can select the columns 'drive-wheels', 'body-style' and 'price', then assign it to the variable "cars\_group\_one".

```
In [91]: 1 cars_group_one = cars[['drive-wheels', 'body-style', 'price']]
```

We can then calculate the average price for each of the different categories of data.

```
In [92]: 1 # grouping results
2 cars_group_one = cars_group_one.groupby(['drive-wheels'], as_index=False).mean()
3 cars_group_one
```

```
Out[92]:
```

	drive-wheels	price
0	4wd	10241.000000
1	fwd	9244.779661
2	rwd	19757.613333

From our data, it seems rear-wheel drive vehicles are, on average, the most expensive, while 4-wheel and front-wheel are approximately the same in price.

let's group by both 'drive-wheels' and 'body-style'. This groups the dataframe by the unique combinations 'drive-wheels' and 'body-style'. We can store the results in the variable 'grouped\_test1'.

```
In [93]: 1 # grouping results
2 df_gptest = cars[['drive-wheels', 'body-style', 'price']]
3 grouped_test1 = df_gptest.groupby(['drive-wheels', 'body-style'], as_index=False)
4 grouped_test1
```

```
Out[93]:
```

	drive-wheels	body-style	price
0	4wd	hatchback	7603.000000
1	4wd	sedan	12647.333333
2	4wd	wagon	9095.750000
3	fwd	convertible	11595.000000
4	fwd	hardtop	8249.000000
5	fwd	hatchback	8396.387755
6	fwd	sedan	9811.800000
7	fwd	wagon	9997.333333
8	rwd	convertible	23949.600000
9	rwd	hardtop	24202.714286
10	rwd	hatchback	14337.777778
11	rwd	sedan	21711.833333
12	rwd	wagon	16994.222222

This grouped data is much easier to visualize when it is made into a pivot table. A pivot table is like an Excel spreadsheet, with one variable along the column and another along the row. We can convert the dataframe to a pivot table using the method "pivot" to create a pivot table from the groups.

In this case, we will leave the drive-wheel variable as the rows of the table, and pivot body-style to become the columns of the table:

```
In [94]: 1 grouped_pivot = grouped_test1.pivot(index='drive-wheels', columns='body-style')
2 grouped_pivot
```

```
Out[94]:
```

		price				
	body-style	convertible	hardtop	hatchback	sedan	wagon
drive-wheels						
	4wd	NaN	NaN	7603.000000	12647.333333	9095.750000
	fwd	11595.0	8249.000000	8396.387755	9811.800000	9997.333333
	rwd	23949.6	24202.714286	14337.777778	21711.833333	16994.222222

Often, we won't have data for some of the pivot cells. We can fill these missing cells with the value 0, but any other value could potentially be used as well.

```
In [95]: 1 grouped_pivot = grouped_pivot.fillna(0) #fill missing values with 0
        2 grouped_pivot
```

```
Out[95]:
```

		price				
	body-style	convertible	hardtop	hatchback	sedan	wagon
<b>drive-wheels</b>						
	<b>4wd</b>	0.0	0.000000	7603.000000	12647.333333	9095.750000
	<b>fwd</b>	11595.0	8249.000000	8396.387755	9811.800000	9997.333333
	<b>rwd</b>	23949.6	24202.714286	14337.777778	21711.833333	16994.222222

```
In [96]: 1 df_gptest2 = cars[['body-style', 'price']]
        2 grouped_test_bodystyle = df_gptest2.groupby(['body-style'], as_index= False).m
        3 grouped_test_bodystyle
```


```
Out[96]:
```

	body-style	price
0	convertible	21890.500000
1	hardtop	22208.500000
2	hatchback	9957.441176
3	sedan	14459.755319
4	wagon	12371.960000

### Variables: Drive Wheels and Body Style vs Price

Let's use a heat map to visualize the relationship between Body Style vs Price.

```
In [97]: 1 #use the grouped results
        2 plt.pcolor(grouped_pivot, cmap='RdBu')
        3 plt.colorbar()
        4 plt.show()
```




The heatmap plots the target variable (price) proportional to colour with respect to the variables 'drive-wheel' and 'body-style' in the vertical and horizontal axis respectively. This allows us to

visualize how the price is related to 'drive-wheel' and 'body-style'.

The default labels convey no useful information to us. Let's change that:

```
In [98]: 1 fig, ax = plt.subplots()
2 im = ax.pcolor(grouped_pivot, cmap='RdBu')
3
4 #Label names
5 row_labels = grouped_pivot.columns.levels[1]
6 col_labels = grouped_pivot.index
7
8 #move ticks and labels to the center
9 ax.set_xticks(np.arange(grouped_pivot.shape[1]) + 0.5, minor=False)
10 ax.set_yticks(np.arange(grouped_pivot.shape[0]) + 0.5, minor=False)
11
12 #insert labels
13 ax.set_xticklabels(row_labels, minor=False)
14 ax.set_yticklabels(col_labels, minor=False)
15
16 #rotate label if too long
17 plt.xticks(rotation=90)
18
19 fig.colorbar(im)
20 plt.show()
```



## Correlation and Causation

**Correlation:** a measure of the extent of interdependence between variables.

**Causation:** the relationship between cause and effect between two variables.

It is important to know the difference between these two and that correlation does not imply causation. Determining correlation is much simpler the determining causation as causation may require independent experimentation.

## Pearson Correlation

The Pearson Correlation measures the linear dependence between two variables X and Y.

The resulting coefficient is a value between -1 and 1 inclusive, where:

- **1**: Total positive linear correlation.
- **0**: No linear correlation, the two variables most likely do not affect each other.
- **-1**: Total negative linear correlation.

Pearson Correlation is the default method of the function "corr". Like before we can calculate the Pearson Correlation of the of the 'int64' or 'float64' variables.

In [99]: 1 cars.corr()

Out[99]:

	symboling	normalized-losses	wheel-base	length	width	height	curb-weight	engine-size
symboling	1.000000	0.466264	-0.535987	-0.365404	-0.242423	-0.550160	-0.233118	-0.110581
normalized-losses	0.466264	1.000000	-0.056661	0.019424	0.086802	-0.373737	0.099404	0.112360
wheel-base	-0.535987	-0.056661	1.000000	0.876024	0.814507	0.590742	0.782097	0.572027
length	-0.365404	0.019424	0.876024	1.000000	0.857170	0.492063	0.880665	0.685025
width	-0.242423	0.086802	0.814507	0.857170	1.000000	0.306002	0.866201	0.729436
height	-0.550160	-0.373737	0.590742	0.492063	0.306002	1.000000	0.307581	0.074694
curb-weight	-0.233118	0.099404	0.782097	0.880665	0.866201	0.307581	1.000000	0.849072
engine-size	-0.110581	0.112360	0.572027	0.685025	0.729436	0.074694	0.849072	1.000000
bore	-0.140019	-0.029862	0.493244	0.608971	0.544885	0.180449	0.644060	0.572027
stroke	-0.000059	0.059131	0.155225	0.121904	0.188301	-0.068633	0.166038	0.199436
compression-ratio	-0.182196	-0.114713	0.250313	0.159733	0.189867	0.259737	0.156433	0.028027
horsepower	0.075810	0.217300	0.371178	0.579795	0.615056	-0.087001	0.757981	0.822027
peak-rpm	0.279740	0.239543	-0.360305	-0.285970	-0.245800	-0.309974	-0.279361	-0.256027
highway-L/100km	-0.029807	0.181189	0.577576	0.707108	0.736728	0.084301	0.836921	0.783027
price	-0.082391	0.133999	0.584642	0.690628	0.751265	0.135486	0.834415	0.872027
city-L/100km	0.066171	0.238567	0.476153	0.657373	0.673363	0.003811	0.785353	0.745027
fuel-type-diesel	-0.196735	-0.101546	0.307237	0.211187	0.244356	0.281578	0.221046	0.070027
fuel-type-gas	0.196735	0.101546	-0.307237	-0.211187	-0.244356	-0.281578	-0.221046	-0.070027
aspiration-std	0.054615	0.006911	-0.256889	-0.230085	-0.305732	-0.090336	-0.321955	-0.110581
aspiration-turbo	-0.054615	-0.006911	0.256889	0.230085	0.305732	0.090336	0.321955	0.110581



sometimes we would like to know the significant of the correlation estimate.

### P-value:

What is this P-value? The P-value is the probability value that the correlation between these two variables is statistically significant. Normally, we choose a significance level of 0.05, which means that we are 95% confident that the correlation between the variables is significant.

By convention, when the

- p-value is  $< 0.001$ : we say there is strong evidence that the correlation is significant.
- the p-value is  $< 0.05$ : there is moderate evidence that the correlation is significant.
- the p-value is  $< 0.1$ : there is weak evidence that the correlation is significant.
- the p-value is  $> 0.1$ : there is no evidence that the correlation is significant.

We can obtain this information using "stats" module in the "scipy" library.

```
In [100]: 1 from scipy import stats
```

## Wheel-base vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'wheel-base' and 'price'.

```
In [101]: 1 pearson_coef, p_value = stats.pearsonr(cars['wheel-base'], cars['price'])
          2 print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value
```

```
The Pearson Correlation Coefficient is 0.5846418222655081 with a P-value of P
= 8.076488270732955e-20
```

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between wheel-base and price is statistically significant, although the linear relationship isn't extremely strong (~0.585)

## Horsepower vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'horsepower' and 'price'.

```
In [102]: 1 pearson_coef, p_value = stats.pearsonr(cars['horsepower'], cars['price'])
          2 print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value
```

```
The Pearson Correlation Coefficient is 0.8096068016571052 with a P-value of P
= 6.273536270651004e-48
```

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between horsepower and price is statistically significant, and the linear relationship is quite strong ( $\sim 0.809$ , close to 1)

## Length vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'length' and 'price'.

```
In [103]: 1 pearson_coef, p_value = stats.pearsonr(cars['length'], cars['price'])
          2 print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P")
```

The Pearson Correlation Coefficient is 0.6906283804483642 with a P-value of P = 8.016477466158713e-30

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between length and price is statistically significant, and the linear relationship is moderately strong ( $\sim 0.691$ ).

## Width vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'width' and 'price':

```
In [104]: 1 pearson_coef, p_value = stats.pearsonr(cars['width'], cars['price'])
          2 print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P")
```

The Pearson Correlation Coefficient is 0.7512653440522673 with a P-value of P = 9.20033551048166e-38

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between width and price is statistically significant, and the linear relationship is quite strong ( $\sim 0.751$ ).

## Curb-weight vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'curb-weight' and 'price':

```
In [105]: 1 pearson_coef, p_value = stats.pearsonr(cars['curb-weight'], cars['price'])
          2 print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P")
```

The Pearson Correlation Coefficient is 0.8344145257702846 with a P-value of P = 2.1895772388936997e-53

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between curb-weight and price is statistically significant, and the linear relationship is quite strong ( $\sim 0.834$ ).

## Engine-size vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'engine-size' and 'price':

```
In [106]: 1 pearson_coef, p_value = stats.pearsonr(cars['engine-size'], cars['price'])
          2 print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P
```

```
The Pearson Correlation Coefficient is 0.8723351674455185 with a P-value of P
= 9.265491622197996e-64
```

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between engine-size and price is statistically significant, and the linear relationship is very strong ( $\sim 0.872$ ).

## Bore vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'bore' and 'price':

```
In [107]: 1 pearson_coef, p_value = stats.pearsonr(cars['bore'], cars['price'])
          2 print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P
```

```
The Pearson Correlation Coefficient is 0.5431553832626602 with a P-value of P
= 8.049189483935364e-17
```

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between bore and price is statistically significant, but the linear relationship is only moderate ( $\sim 0.521$ ).

## city-L/100km vs Price

```
In [108]: 1 pearson_coef, p_value = stats.pearsonr(cars['city-L/100km'], cars['price'])
          2 print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P
```

```
The Pearson Correlation Coefficient is 0.7898975136626942 with a P-value of P
= 3.9031064009399405e-44
```

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between city-L/100km and price is statistically significant, and the coefficient of  $\sim 0.789$  shows that the relationship is moderately strong.

## Highway-L/100km vs Price

```
In [109]: 1 pearson_coef, p_value = stats.pearsonr(cars['highway-L/100km'], cars['price'])
          2 print( "The Pearson Correlation Coefficient is", pearson_coef, " with a P-val
```

The Pearson Correlation Coefficient is 0.8011176263981975 with a P-value of P = 3.0467845810412534e-46

### **Conclusion:**

Since the p-value is  $< 0.001$ , the correlation between highway-L/100km and price is statistically significant, and the coefficient of  $\sim 0.801$  shows that the relationship is moderately strong.

## Conclusion: Important Variables

We now have a better idea of what our data looks like and which variables are important to take into account when predicting the car price. We have narrowed it down to the following variables:

Continuous numerical variables:

- Length
- Width
- Curb-weight
- Engine-size
- Horsepower
- City-L/100km
- Highway-L/100km
- Wheel-base
- Bore

Categorical variables:

- Drive-wheels

As we now move into building machine learning models to automate our analysis, feeding the model with variables that meaningfully affect our target variable will improve our model's prediction performance.

## Model Development

we will develop several models that will predict the price of the car using the variables or features. This is just an estimate but should give us an objective idea of how much the car should cost.

### Linear Regression

### Lets load the modules for linear regression

```
In [183]: 1 from sklearn.linear_model import LinearRegression
```

### Create the linear regression object

```
In [184]: 1 lm = LinearRegression()  
2 lm
```

```
Out[184]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

### How could Highway-L/100km help us predict car price?

we want to look at how highway-mpg can help us predict car price. Using simple linear regression, we will create a linear function with "highway-mpg" as the predictor variable and the "price" as the response variable.

```
In [185]: 1 X = cars[['highway-L/100km']]  
2 Y = cars['price']
```

Fit the linear model using highway-L/100km.

```
In [186]: 1 lm.fit(X,Y)
```

```
Out[186]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

We can output a prediction

```
In [114]: 1 Yhat=lm.predict(X)  
2 Yhat[0:5]
```

```
Out[114]: array([15485.52737455, 15485.52737455, 16643.34931414, 12475.19033163,  
                22327.2024721 ])
```

```
In [187]: 1 lm.score(X,Y)
```

```
Out[187]: 0.6417894513258818
```

### What is the value of the intercept (a)?

```
In [115]: 1 lm.intercept_
```

```
Out[115]: -14617.843054664598
```

### What is the value of the Slope (b)?

```
In [116]: 1 lm.coef_
```

```
Out[116]: array([3458.68511314])
```

## What is the final estimated linear model we get?

As we saw above, we should get a final linear model with the structure:

$$\hat{Y} = a + bX$$


Plugging in the actual values we get:

**price** = -14617.84 + 3458.68 x **highway-L/100km**

## Model Evaluation using Visualization


```
In [117]: 1 width = 12
          2 height = 10
          3 plt.figure(figsize=(width, height))
          4 sns.regplot(x="highway-L/100km", y="price", data=cars)
          5 plt.ylim(0,)
```

Out[117]: (0, 47827.10608389057)



We can see from this plot that price is positively correlated to highway-L/100km, since the regression slope is positive. Since data is too far off from the line, this linear model might not be the best model for this data.


```
In [118]: 1 width = 12
2 height = 10
3 plt.figure(figsize=(width, height))
4 sns.residplot(cars['highway-L/100km'], cars['price'])
5 plt.show()
```



We can see from this residual plot that the residuals are not randomly spread around the x-axis, which leads us to believe that maybe a non-linear model is more appropriate for this data.



```
In [119]: 1 plt.figure(figsize=(width, height))
2
3
4 ax1 = sns.distplot(cars['price'], hist=False, color="r", label="Actual Value"
5 sns.distplot(Yhat, hist=False, color="b", label="Fitted Values" , ax=ax1)
6
7
8 plt.title('Actual vs Fitted Values for Price')
9 plt.xlabel('Price (in dollars)')
10 plt.ylabel('Proportion of Cars')
11
12 plt.show()
13 plt.close()
```



We can see that the fitted values are reasonably close to the actual values, since the two distributions overlap a bit. However, there is definitely some room for improvement.

**How could city-L/100km help us predict car price?**

```
In [120]: 1 lm1 = LinearRegression()  
          2 lm1  
          3
```

```
Out[120]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

we want to look at how city-L/100km can help us predict car price. Using simple linear regression, we will create a linear function with "highway-mpg" as the predictor variable and the "price" as the response variable.

```
In [121]: 1 X = cars[['city-L/100km']]  
          2 Y = cars['price']
```

Fit the linear model using highway-L/100km.

```
In [122]: 1 lm1.fit(X,Y)
```

```
Out[122]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

We can output a prediction

```
In [123]: 1 Yhat=lm1.predict(X)  
          2 Yhat[0:5]
```

```
Out[123]: array([16293.8802044 , 16293.8802044 , 19211.26183196, 12829.48952166,  
                  20913.06778137])
```

```
In [124]: 1 print("intercept:",lm1.intercept_,"slope:",lm1.coef_)
```

```
intercept: -11421.245257455565 slope: [2476.67078595]
```

Estimated Linear model:

$$\hat{Y} = a + bX$$


Plugging in the actual values we get:

**price** = -11421.24 + 2476.67 x **city-L/100km**

## Model Evaluation using Visualization


```
In [125]: 1 width = 12
2 height = 10
3 plt.figure(figsize=(width, height))
4 sns.regplot(x="city-L/100km", y="price", data=cars)
5 plt.ylim(0,)
6 plt.title("Regression Plot")
```

```
Out[125]: Text(0.5, 1.0, 'Regression Plot')
```




We can see from this plot that price is positively correlated to city-L/100km, since the regression slope is positive. Since data is too far off from the line, this linear model might not be the best model for this data.

```
In [126]: 1 width = 12
2 height = 10
3 plt.figure(figsize=(width, height))
4 sns.residplot(cars['city-L/100km'], cars['price'])
5 plt.title("residual Plot")
6 plt.show()
```



We can see from this residual plot that the residuals are not randomly spread around the x-axis, which leads us to believe that maybe a non-linear model is more appropriate for this data.

```
In [127]: 1 plt.figure(figsize=(width, height))
2
3
4 ax1 = sns.distplot(cars['price'], hist=False, color="r", label="Actual Value"
5 sns.distplot(Yhat, hist=False, color="b", label="Fitted Values" , ax=ax1)
6
7
8 plt.title('Actual vs Fitted Values for Price')
9 plt.xlabel('Price (in dollars)')
10 plt.ylabel('Proportion of Cars')
11
12 plt.show()
13 plt.close()
```



We can see that the fitted values are reasonably close to the actual values, since the two distributions overlap a bit. However, there is definitely some room for improvement.

## How could Engine-size help us predict car price?

```
In [172]: 1 lm2 = LinearRegression()
          2 lm2
```

```
Out[172]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [173]: 1 X = cars[['engine-size']]
          2 Y = cars['price']
```

```
In [174]: 1 lm2.fit(X,Y)
```

```
Out[174]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [175]: 1 yhat = lm2.predict(X)
          2 yhat[0:5]
```

```
Out[175]: array([13728.4631336 , 13728.4631336 , 17399.38347881, 10224.40280408,
                  14729.62322775])
```

```
In [176]: 1 print("intercept:",lm2.intercept_,"slope:",lm2.coef_)
```

```
intercept: -7963.338906281049 slope: [166.86001569]
```

Estimated Linear model:

$$\hat{Y} = a + bX$$


Plugging in the actual values we get:

**price** = -7963.33 + 166.86 x **engine-size**

## Model Evaluation using Visualization


```
In [134]: 1 width = 12
2 height = 10
3 plt.figure(figsize=(width, height))
4 sns.regplot(x="engine-size", y="price", data=cars)
5 plt.ylim(0,)
6 plt.title("Regression Plot")
```

Out[134]: Text(0.5, 1.0, 'Regression Plot')



We can see from this plot that price is positively correlated to engine-size, since the regression slope is positive. Since data is too far off from the line, this linear model might not be the best model for this data.


```
In [135]: 1 width = 12
2 height = 10
3 plt.figure(figsize=(width, height))
4 sns.residplot(cars['engine-size'], cars['price'])
5 plt.title("residual Plot")
6 plt.show()
```



We can see from this residual plot that the residuals are not randomly spread around the x-axis, which leads us to believe that maybe a non-linear model is more appropriate for this data.



```
In [136]: 1 plt.figure(figsize=(width, height))
2
3
4 ax1 = sns.distplot(cars['price'], hist=False, color="r", label="Actual Value"
5 sns.distplot(yhat, hist=False, color="b", label="Fitted Values" , ax=ax1)
6
7
8 plt.title('Actual vs Fitted Values for Price')
9 plt.xlabel('Price (in dollars)')
10 plt.ylabel('Proportion of Cars')
11
12 plt.show()
13 plt.close()
```



We can see that the fitted values are reasonably close to the actual values, since the two distributions overlap a lot. However, there is definitely some room for improvement.

## Multiple Linear Regression

In [137]: 1 cars.corr()

Out[137]:

	symboling	normalized-losses	wheel-base	length	width	height	curb-weight	engine-size
symboling	1.000000	0.466264	-0.535987	-0.365404	-0.242423	-0.550160	-0.233118	-0.110581
normalized-losses	0.466264	1.000000	-0.056661	0.019424	0.086802	-0.373737	0.099404	0.112360
wheel-base	-0.535987	-0.056661	1.000000	0.876024	0.814507	0.590742	0.782097	0.572027
length	-0.365404	0.019424	0.876024	1.000000	0.857170	0.492063	0.880665	0.685025
width	-0.242423	0.086802	0.814507	0.857170	1.000000	0.306002	0.866201	0.729436
height	-0.550160	-0.373737	0.590742	0.492063	0.306002	1.000000	0.307581	0.074694
curb-weight	-0.233118	0.099404	0.782097	0.880665	0.866201	0.307581	1.000000	0.849072
engine-size	-0.110581	0.112360	0.572027	0.685025	0.729436	0.074694	0.849072	1.000000
bore	-0.140019	-0.029862	0.493244	0.608971	0.544885	0.180449	0.644060	0.572027
stroke	-0.000059	0.059131	0.155225	0.121904	0.188301	-0.068633	0.166038	0.155225

Length Width Curb-weight Engine-size Horsepower City-L/100km Highway-L/100km Wheel-base Bore,,,,,,Let's develop a model using these variables as the predictor variables.

In [195]: 1 z = cars[["curb-weight", "engine-size", "horsepower", "highway-L/100km"]]

Fit the linear model using the four above-mentioned variables.

In [196]: 1 lm\_m = LinearRegression()  
2 lm\_m

Out[196]: LinearRegression(copy\_X=True, fit\_intercept=True, n\_jobs=None, normalize=False)

In [197]: 1 lm\_m.fit(z, cars["price"])

Out[197]: LinearRegression(copy\_X=True, fit\_intercept=True, n\_jobs=None, normalize=False)

In [198]: 1 yhat = lm\_m.predict(z)  
2 yhat[0:5]

Out[198]: array([14055.08612634, 14055.08612634, 18638.8884263 , 10758.2803759 ,  
16670.07025864])


```
In [199]: 1 print("intercept:",lm_m.intercept_,"slope:",lm_m.coef_)
```

```
intercept: -14385.634549360078 slope: [ 3.50038215 85.37370862 36.6377371 5
00.51979785]
```

```
In [200]: 1 lm_m.score(z,cars["price"])
```

```
Out[200]: 0.811811561534475
```

```
In [143]: 1 plt.figure(figsize=(width, height))
2
3
4 ax1 = sns.distplot(cars['price'], hist=False, color="r", label="Actual Value"
5 sns.distplot(yhat, hist=False, color="b", label="Fitted Values" , ax=ax1)
6
7
8 plt.title('Actual vs Fitted Values for Price')
9 plt.xlabel('Price (in dollars)')
10 plt.ylabel('Proportion of Cars')
11
12 plt.show()
13 plt.close()
```



We can see that the fitted values are reasonably close to the actual values, since the two distributions overlap a lot. However, there is definitely some room for improvement.

```
In [144]: 1 lm_m.score(z, cars["price"])
```

```
Out[144]: 0.811811561534475
```

## Polynomial Regression

highway-L/100km(corr() with price = 0.801)

```
In [205]: 1 x = cars["highway-L/100km"]
          2 y = cars["price"]
          3
```


Let's fit the polynomial using the function **polyfit**, then use the function **poly1d** to display the polynomial function.

```
In [206]: 1 #We are using a polynomial of degree 2.
          2 f = np.polyfit(x,y,2)
          3 p = np.poly1d(f)
          4 print(p)
```

```
2
266.1 x - 1234 x + 5019
```

```
In [215]: 1 def PlotPolly(model, independent_variable, dependent_variabble, Name):
          2     x_new = np.linspace(0, 20, 10)
          3     y_new = model(x_new)
          4
          5     plt.plot(independent_variable, dependent_variabble, '.', x_new, y_new, '-')
          6     plt.title('Polynomial Fit with Matplotlib for Price ~ Length')
          7     ax = plt.gca()
          8     ax.set_facecolor((0.898, 0.898, 0.898))
          9     fig = plt.gcf()
         10     plt.xlabel(Name)
         11     plt.ylabel('Price of Cars')
         12
         13     plt.show()
         14     plt.close()
```

```
In [216]: 1 PlotPolly(p, x, y, 'highway-L/100km')
```



```
In [217]: 1 from sklearn.metrics import r2_score
```

```
In [218]: 1 r_squared = r2_score(y, p(x))  
2 print('The R-square value is: ', r_squared)
```

The R-square value is: 0.6733365470678851

```
In [219]: 1 from sklearn.metrics import mean_squared_error  
2 mean_squared_error(cars['price'], p(x))
```

Out[219]: 20528072.06493496

Performing a polynomial transform on multiple features:(Z)

```
In [150]: 1 from sklearn.preprocessing import PolynomialFeatures
```

We create a **PolynomialFeatures** object of degree 2:

```
In [151]: 1 pr=PolynomialFeatures(degree=2)
          2 pr
```

```
Out[151]: PolynomialFeatures(degree=2, include_bias=True, interaction_only=False,
                             order='C')
```

```
In [152]: 1 Z_pr=pr.fit_transform(z)
```

The original data is of 201 samples and 4 features

```
In [153]: 1 z.shape
```

```
Out[153]: (201, 4)
```

```
In [154]: 1 Z_pr.shape
```

```
Out[154]: (201, 15)
```

after the transformation, there 201 samples and 15 features

## Pipeline

```
In [155]: 1 from sklearn.pipeline import Pipeline
          2 from sklearn.preprocessing import StandardScaler
```

```
In [156]: 1 Input=[('scale',StandardScaler()), ('polynomial', PolynomialFeatures(include_
```

```
In [157]: 1 pipe=Pipeline(Input)
          2 pipe
```

```
Out[157]: Pipeline(memory=None,
                  steps=[('scale',
                          StandardScaler(copy=True, with_mean=True, with_std=True)),
                          ('polynomial',
                           PolynomialFeatures(degree=2, include_bias=False,
                                                interaction_only=False, order='C')),
                          ('model',
                           LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                                                normalize=False))],
                  verbose=False)
```

We can normalize the data, perform a transform and fit the model simultaneously.

```
In [158]: 1 pipe.fit(z,cars["price"])
```

```
Out[158]: Pipeline(memory=None,
                  steps=[('scale',
                          StandardScaler(copy=True, with_mean=True, with_std=True)),
                          ('polynomial',
                          PolynomialFeatures(degree=2, include_bias=False,
                                              interaction_only=False, order='C')),
                          ('model',
                          LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                                              normalize=False))],
                  verbose=False)
```

Similarly, we can normalize the data, perform a transform and produce a prediction simultaneously

```
In [159]: 1 ypipe=pipe.predict(z)
          2 ypipe[0:4]
```

```
Out[159]: array([12395.58706934, 12395.58706934, 18422.45365339,  9979.7757077 ])
```

```
In [160]: 1 pipe.score(z,cars["price"])
```

```
Out[160]: 0.8531061720624729
```

## Measures for In-Sample Evaluation

When evaluating our models, not only do we want to visualize the results, but we also want a quantitative measure to determine how accurate the model is.

Two very important measures that are often used in Statistics to determine the accuracy of a model are:

- **R<sup>2</sup> / R-squared**
- **Mean Squared Error (MSE)**

### R-squared

R squared, also known as the coefficient of determination, is a measure to indicate how close the data is to the fitted regression line.

The value of the R-squared is the percentage of variation of the response variable (y) that is explained by a linear model.

### Mean Squared Error (MSE)

The Mean Squared Error measures the average of the squares of errors, that is, the difference between actual value (y) and the estimated value ( $\hat{y}$ ).

## Model 1: Simple Linear Regression

Let's calculate the R<sup>2</sup>

```
In [189]: 1 #highway-L/100km_fit
          2 lm.fit(X, Y)
          3 # Find the R^2
          4 print('The R-square is: ', lm.score(X, Y))
```

The R-square is: 0.6417894513258818

We can say that ~ 64.178% of the variation of the price is explained by this simple linear model "highway-L/100km\_fit".

Let's calculate the MSE

```
In [190]: 1 Yhat=lm.predict(X)
          2 print('The output of the first four predicted value is: ', Yhat[0:4])
```

The output of the first four predicted value is: [15485.52737455 15485.52737455 16643.34931414 12475.19033163]

we compare the predicted results with the actual results

```
In [192]: 1 mse = mean_squared_error(cars['price'], Yhat)
          2 print('The mean square error of price and predicted value is: ', mse)
```

The mean square error of price and predicted value is: 22510543.777085222

## Model 2: Multiple Linear Regression

Let's calculate the R^2

```
In [202]: 1 # fit the model
          2 lm_m.fit(z, cars['price'])
          3 # Find the R^2
          4 print('The R-square is: ', lm.score(z, cars['price']))
```

The R-square is: 0.811811561534475

We can say that ~ 81.181 % of the variation of price is explained by this multiple linear regression "multi\_fit".

Let's calculate the MSE

```
In [203]: 1 Y_predict_multifit = lm_m.predict(z)
```



```
In [204]: 1 print('The mean square error of price and predicted value using multfit is:
          2 mean_squared_error(cars['price'], Y_predict_multifit))
```

The mean square error of price and predicted value using multfit is: 11826072.956532082

## Model 3: Polynomial Fit

Let's calculate the  $R^2$

```
In [220]: 1 r_squared = r2_score(y, p(x))
          2 print('The R-square value is: ', r_squared)
```

The R-square value is: 0.6733365470678851

We can say that ~ 67.333 % of the variation of price is explained by this polynomial fit

We can also calculate the MSE:

```
In [221]: 1 mean_squared_error(cars["price"], p(x))
```

Out[221]: 20528072.06493496

## Decision Making: Determining a Good Model Fit

Now that we have visualized the different models, and generated the R-squared and MSE values for the fits, how do we determine a good model fit?

- *What is a good R-squared value?*

When comparing models, **the model with the higher R-squared value is a better fit** for the data.

- *What is a good MSE?*

When comparing models, **the model with the smallest MSE value is a better fit** for the data.

**Let's take a look at the values for the different models.**

Simple Linear Regression: Using Highway-L/100km as a Predictor Variable of Price.

- R-squared: 0.64
- MSE: 22510543.77

Multiple Linear Regression: Using Horsepower, Curb-weight, Engine-size, and Highway-L/100km as Predictor Variables of Price.

- R-squared: 0.81
- MSE: 11826072.95

Polynomial Fit: Using Highway-mpg as a Predictor Variable of Price.

- R-squared: 0.67
- MSE: 20528072.06

## Simple Linear Regression model (SLR) vs Multiple Linear Regression model (MLR)

To be able to compare the results of the MLR vs SLR models, we look at a combination of both the R-squared and MSE to make the best conclusion about the fit of the model.

- **MSE** The MSE of SLR is 22510543.77 while MLR has an MSE of 11826072.95. The MSE of MLR is much smaller.
- **R-squared**: In this case, we can also see that there is a big difference between the R-squared of the SLR and the R-squared of the MLR. The R-squared for the SLR (~0.64) is very small compared to the R-squared for the MLR (~0.81).

This R-squared in combination with the MSE show that MLR seems like the better model fit in this case, compared to SLR.

## Simple Linear Model (SLR) vs Polynomial Fit

- **MSE**: We can see that Polynomial Fit brought down the MSE, since this MSE is smaller than the one from the SLR.
- **R-squared**: The R-squared for the Polyfit is larger than the R-squared for the SLR, so the Polynomial Fit also brought up the R-squared quite a bit.

Since the Polynomial Fit resulted in a lower MSE and a higher R-squared, we can conclude that this was a better fit model than the simple linear regression for predicting Price with Highway-mpg as a predictor variable.

## Multiple Linear Regression (MLR) vs Polynomial Fit

- **MSE**: The MSE for the MLR is smaller than the MSE for the Polynomial Fit.
- **R-squared**: The R-squared for the MLR is also much larger than for the Polynomial Fit.

## Conclusion:

Comparing these three models, we conclude that **the MLR model is the best model** to be able to predict price from our dataset. This result makes sense, since we have 27 variables in total, and we know that more than one of those variables are potential predictors of the final car price.

In [ ]:

1

